

logistic regression practise

February 11, 2023

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[5]: df=pd.read_csv('diabetes.csv')
```

```
[6]: df.head()
```

```
[6]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[15]: df.isna().sum()
```

```
[15]: Pregnancies          0
Glucose                  0
BloodPressure            0
SkinThickness            0
Insulin                  0
BMI                      0
DiabetesPedigreeFunction  0
Age                      0
Outcome                  0
dtype: int64
```

```
[52]: df['BMI'].replace([np.inf,-np.inf],np.nan,inplace=True)
```

```
[44]: df['BloodPressure'].unique()
```

```
[44]: array([4.27666612, 4.18965474, 4.15888308, 3.68887945, 4.30406509,
        3.91202301, 4.26694579, 4.24849524, 4.56434819, 4.52178858,
        4.38202663, 4.09434456, 4.4308168 , 3.40119738, 4.47733681,
        4.49980967, 4.54329478, 4.33073334, 4.40671925, 4.31748811,
        4.06044301, 4.35670883, 4.21950771, 4.70048037, 4.02535169,
        4.12713439, 4.44265126, 4.4543473 , 3.87120101, 3.78418963,
        4.17438727, 4.68213123, 4.00733319, 4.80402104, 3.98898405,
        3.95124372, 4.58496748, 4.6443909 , 4.55387689, 3.8286414 ,
        4.62497281, 4.60517019, 4.11087386, 3.17805383, 3.63758616,
        4.66343909, 4.73619845])
```

```
[31]: from sklearn.impute import SimpleImputer
```

```
[32]: si=SimpleImputer()
```

```
[53]: df[['BMI']]=si.fit_transform(df[['BMI']])
```

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
4   Insulin                768 non-null    int64
5   BMI                   768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                   768 non-null    int64
8   Outcome                768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[8]: df.describe()
```

```
[8]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	\
count	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	
std	3.369578	31.972618	19.355807	15.952218	115.244002	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	

50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[11]: plt.figure(figsize=(12,12))
      sns.boxplot(data=df)
```

```
[11]: <AxesSubplot:>
```




```
[14]: df.corr().style.background_gradient()
```

```
[14]: <pandas.io.formats.style.Styler at 0x200ef588fa0>
```


```
[16]: from scipy.stats import skew
```

```
[25]: for i in df[colname]:
        print(i)
        print(skew(df[i]))
        sns.distplot(df[i])
        plt.show()
```

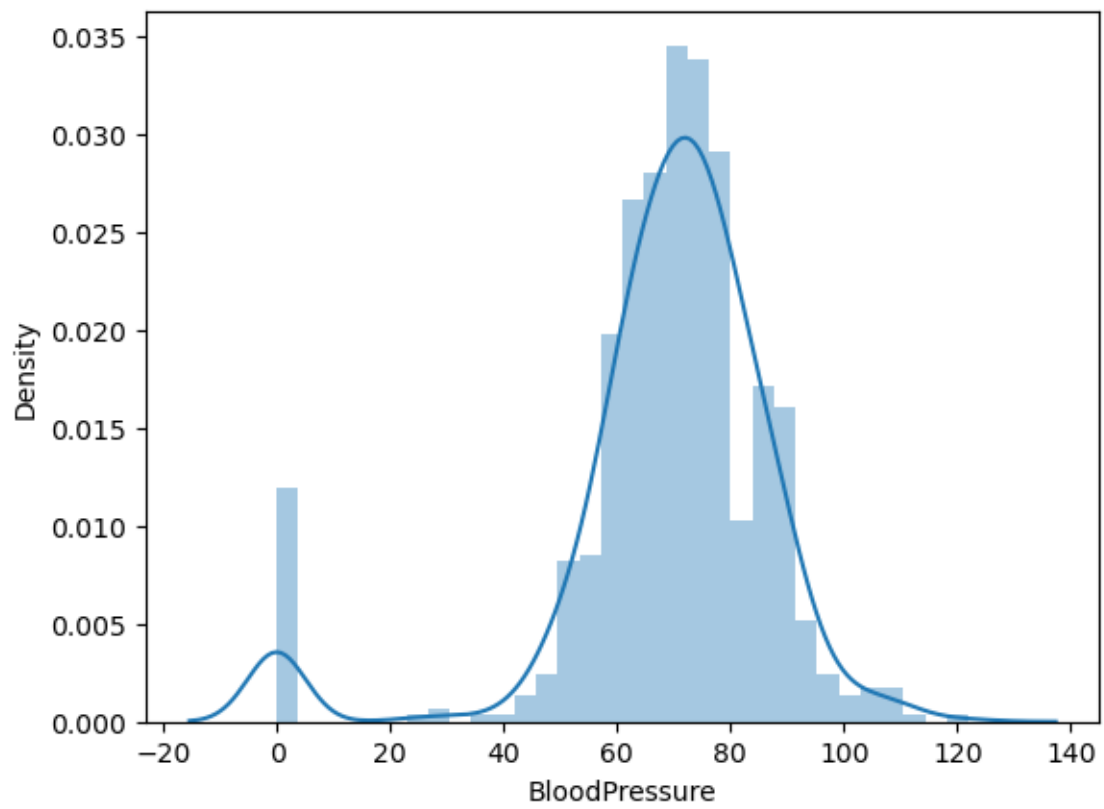
```
Pregnancies
0.8999119408414357
```




Glucose
0.17341395519987735




BloodPressure
-1.8400052311728738




SkinThickness
0.109158762323673




Insulin
2.2678104585131753




BMI
-0.42814327880861786



DiabetesPedigreeFunction
1.9161592037386292




Age
1.127389259531697




```
[26]: df[colname]=np.log(df[colname])
```

```
[54]: for i in df[colname]:  
       print(i)  
       print(skew(df[i]))  
       sns.distplot(df[i])  
       plt.show()
```


Pregnancies
-0.22092086734048239




Glucose
-0.06522741519898129




BloodPressure
-0.8117813252892095




SkinThickness
-0.757853177699361



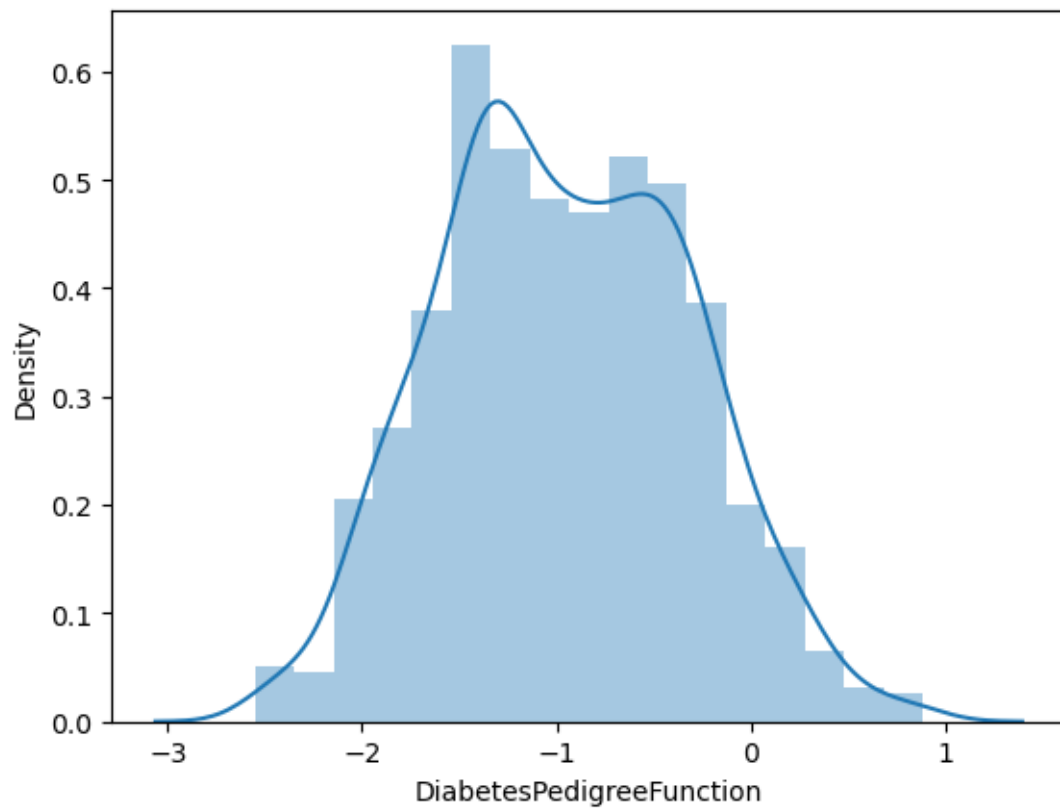
Insulin
-0.1591138900700523




BMI
-0.05242667957133482



DiabetesPedigreeFunction
0.11395456387082803



Age
0.6005702138973051



```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[24]: df[colname]
```

```
[24]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	

763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33
..
763	0.171	63
764	0.340	27
765	0.245	30
766	0.349	47
767	0.315	23

[768 rows x 8 columns]

```
[23]: colname
```

```
[23]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
          'BMI', 'DiabetesPedigreeFunction', 'Age'],
          dtype='object')
```

```
[59]: x=df.iloc[:, :-1]
      y=df.iloc[:, -1]
```


```
[60]: colname=x.columns
```

```
[61]: y
```

```
[61]: 0      1
      1      0
      2      1
      3      0
      4      1
      ..
      763    0
      764    0
      765    0
      766    1
      767    0
      Name: Outcome, Length: 768, dtype: int64
```

```
[55]: plt.figure(figsize=(12,12))
sns.boxplot(data=df)
```

```
[55]: <AxesSubplot:>
```



```
[58]: from sklearn.model_selection import train_test_split
```

```
[63]: xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.3,random_state=1)
```

```
[64]: xtrain
```

```
[64]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
88      2.708050   4.912655    4.248495    3.465736   4.700480  3.613617
```

	DiabetesPedigreeFunction	Age
88	-1.877317	3.761200
467	-0.510826	3.218876
550	-1.589635	3.044522
147	0.336472	3.526361
481	-1.624552	3.367296
..
645	-2.009915	3.401197
715	-0.191161	3.526361
72	-0.539568	3.737670
235	-0.736055	3.258097
37	-0.407968	3.828641

```
[57]: lr=LogisticRegression()
```

```
[65]: lr.fit(xtrain,ytrain)
```

```
[66]: ypred=lr.predict(xtest)
```

```
[67]: ypred
```

23

```
1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0,  
1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0,  
1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,  
0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,  
0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0], dtype=int64)
```

```
[ ]:
```